

Programming Assignment 4 malloc() Replacement

Introduction to Operating Systems
Prof. L-Pin Chang @ NYCU

Objectives

- To replace the original implementation of malloc() and free() with your own version
- To evaluate the performance of Best Fit space allocation algorithm (based on the multilevel list implementation)

malloc()

- Part of the standard C library
- Linux employs the GNU implementation, [glibc](#),

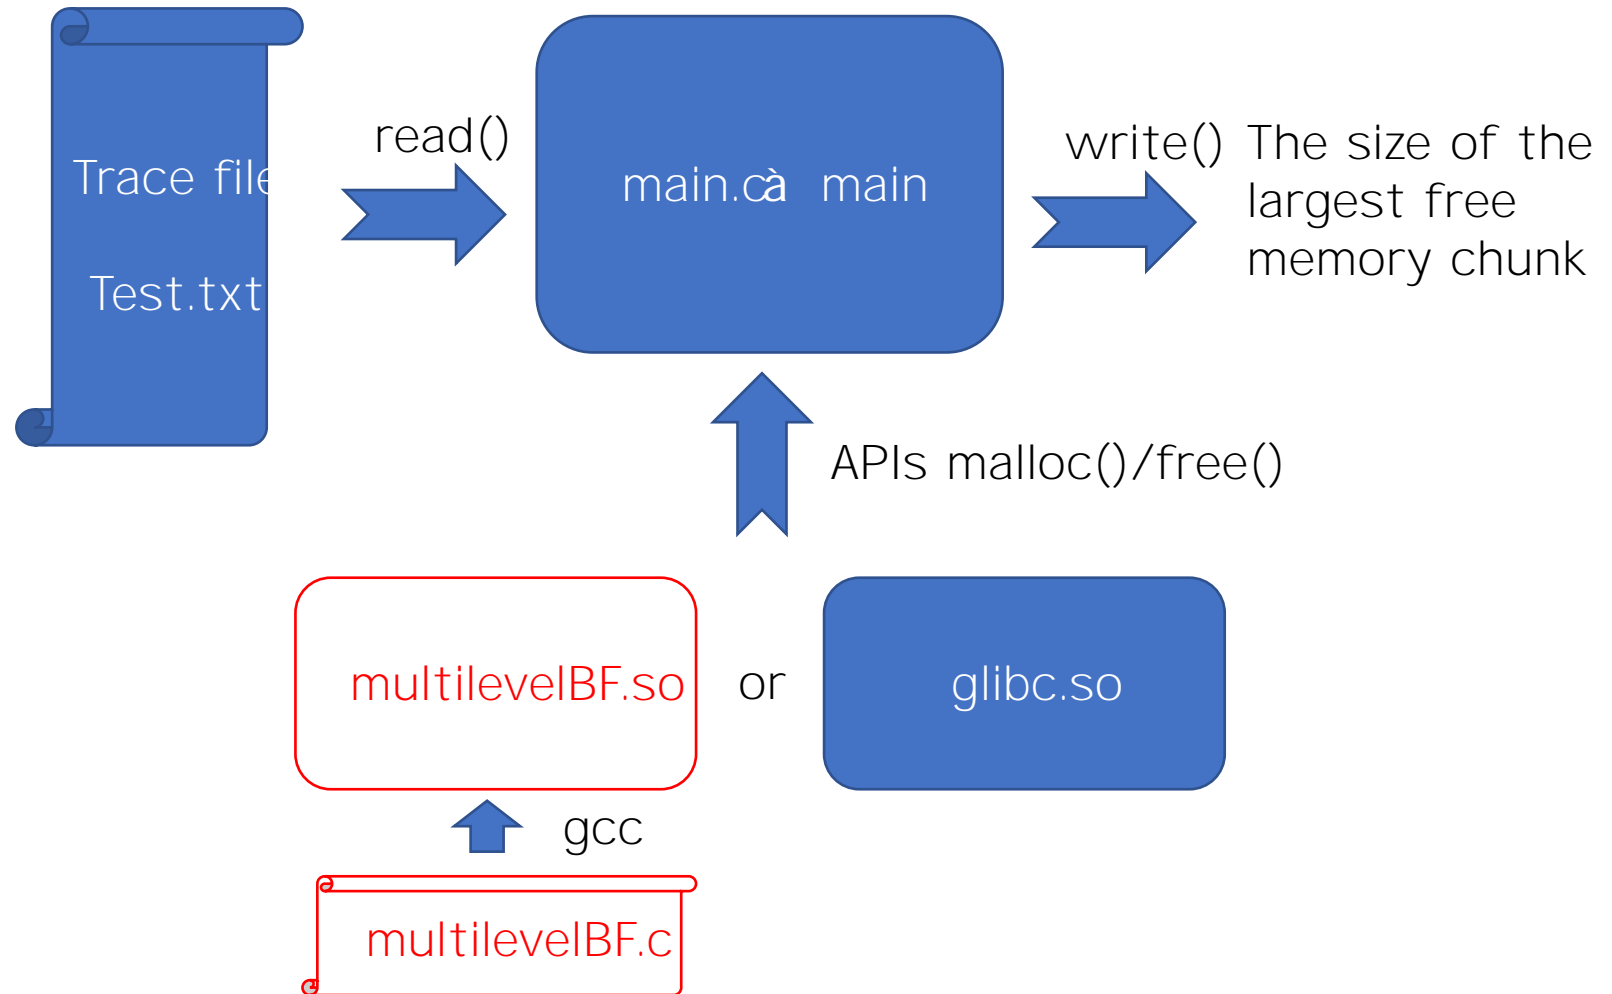
Implementation details about malloc() in

- Small requests ($< \text{M_MMAP_THRESHOLD}$, i.e., 128KB) are serviced using the heap. Heap is resized using `brk()` or `sbrk()` if necessary
- Large requests are serviced by asking the kernel to allocate a piece of anonymous memory using

Assignment Overview

- TA provides two files
 - test.txt: Input file that defines operations of memory allocation and deallocation
 - main.c A program that calls malloc() and free() using the operations in test.txt
- You write one file
 - multilevelBF: your malloc() & free() using Best Fit with multilevel free list

Assignment Overview



Test Flow

- 1) Compile main.c into main and put test.txt in the same d
- 2) Run `./main`
 - Should be no problem
- 3) Compile multilevelBF.c into multilevelBF.so
`$ gcc-shared -fPIC multilevelBF.c -o multilevelBF.so`
- 4) Run `LD_PRELOAD=/path/to/your/multilevelBF.so ./main`
 - Print a result on the screen

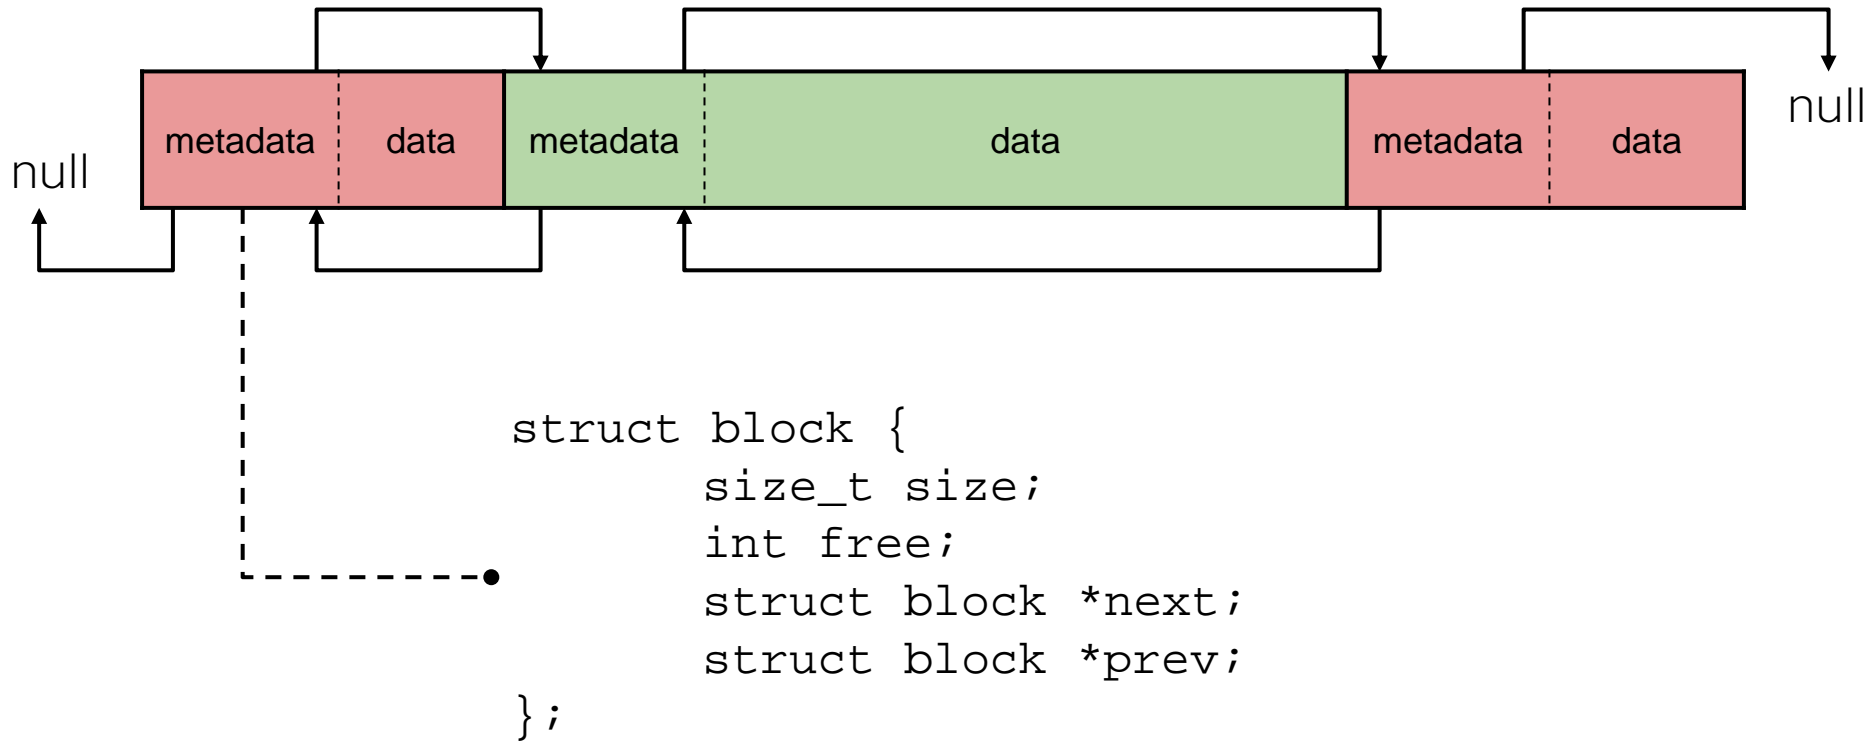
Remark: environment variable: LD_PRELOAD

- A list of additional, user specified, ELF shared objects to be loaded before all others
- malloc() & free() in multilevelBF.so override the original ones

Your Implementation (tilelevelBE)

- On the first malloc()
 - Pre-allocate a memory pool of 20,000 bytes from the kernel using mmap()
 - Initialize metadata for your memory pool
- On subsequent malloc() and free()
 - Process malloc() and free() within the memory pool
- On malloc(0)
 - A fake request that indicates success
 - Print the size of the largest free chunk
 - Call munmap() to release the memory pool

Metadata and Layout

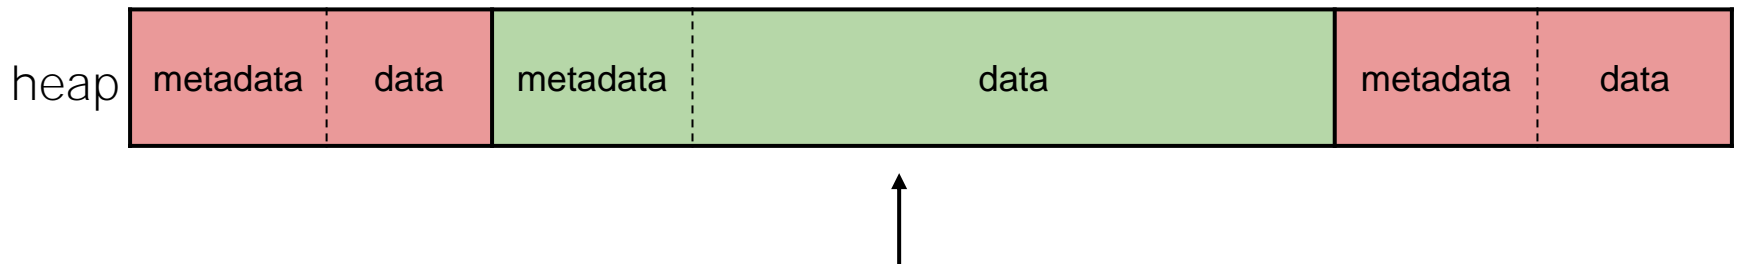


Notice that your header must exactly use 32 bytes.
Use `sizeof()` and padding if necessary.

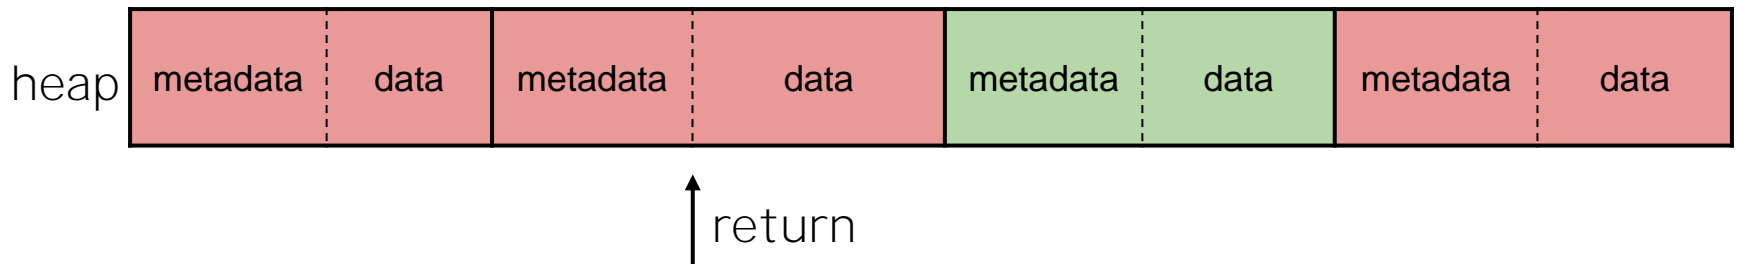
Memory Pool Management

· `void *malloc(size_t size);`

1. choose and split



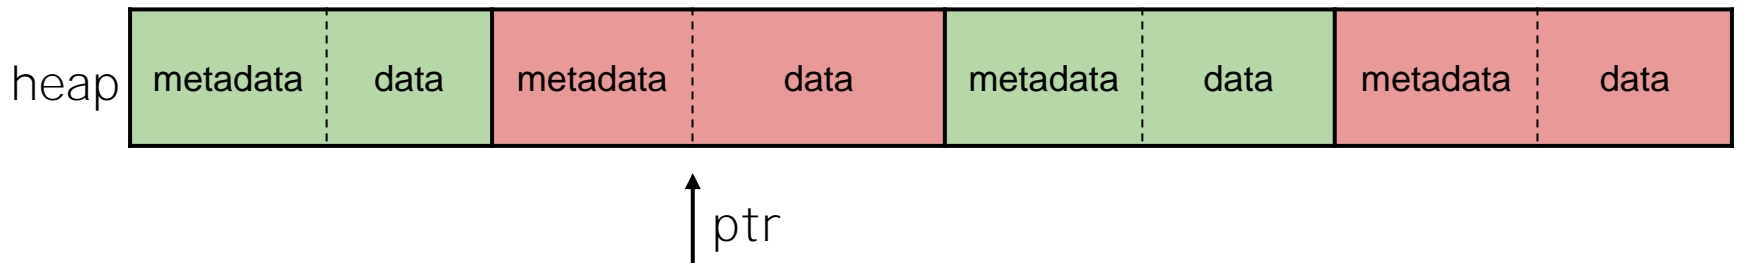
2. return the pointer




Memory Pool Management

· void free(void *ptr);

1. free the memory block




2. merge with free neighbor(s)



Implementation Details (!)

- Chunk list (chunk = space)
 - A list manages all memory chunks, both used and free
 - Initially has only one free memory chunk (20,000 bytes)
- The header of a chunk is exactly 32 bytes
 - Including paddings (if necessary)
- Memory alignment
 - The starting address of the memory pool must be aligned to (this is guaranteed by `mmap()`)
 - The allocation size must be rounded to a multiple of 32
- The memory address returned by `malloc()` must be aligned to 32 bytes for example:
 - The starting memory address of the memory pool is 8192
 - The return address of the first `malloc(31)` is $8192 + 32$
 - The return address of the second `malloc()` is $8192 + 32 + 32 + 32$

Multilevel Free List



- " An array of pointers. A level has a pointer to a list of free blocks whose sizes are 2^i [2
- " On allocation, select the **first best fitting** free block from lists. On release, append the freed block to the **end** of lists.


Details on malloc()

- Use the multilevel free list to find a free block
 - Find the best fitting level (powers of 2)
 - If no free blocks, descend to the next level
 - Each level follows **Best Fit (the first best fitting one)**
 - Get a free block and split if necessary


We assume no header overhead in this example!

Example: `malloc(150)` round to 160

Memory: free/allocated



Free lists





Free lists

Details on free()


- Return a block to the multilevel list
 - If neighbors are also free blocks, merge them
 - Delete old free blocks on merge
 - Append the new free block to the end of a list

Example: free(64)

Memory: free/allocated



Free lists



Free lists

Design Problem

- How do you squeeze the following into a 32B header
 - Free flag
 - Size
 - Two list pointers (pse, vnext) for neighbors
 - Two list pointers (pse, vnext) for free blocks

APIs

- `<sys/mman.h>`
- `mmap()` - creates a new mapping in the virtual address space of the calling process
- `munmap()` - deletes the mappings

- <https://man7.org/linux/man-pages/man2/mmap.2.html>
- <https://man7.org/linux/man-pages/man8/ld.so.8.html>

mmap()

- `void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);`
 - `addr`: NULL for system to choose suitable address
 - `length` : the length of the mapping
 - `prot`: PROT_READ | PROT_WRITE for read and write
 - `flags` MAP_ANON since our mapping is not backed by a file, MAP_PRIVATE let updates invisible to other processes
 - `fd`: -1 for ignored (in conjunction with MAP_ANON)
 - `offset` : 0

munmap()

- `int munmap(void *addr, size_t length);`
 - `addr`: The starting address to be unmapped (must be a multiple of the page size)
 - `length` : the length to be unmapped

Remark: malloc() called within APIs

- You may notice that `main.c` avoids using `fopen()`, `scanf()`, and `printf()` because these APIs call `malloc()` internally and will affect your result (or deadlock your program)
 - `fopen()` -> `open()`
 - `fread()` -> `read()`
 - `fclose()` -> `close()`
- To print out a string
 - Use a local variable string
 - Use `printf()` to format your string
 - Use `wrtie(stdout,*)` to output your string

Input and Output

- Input filename: test.txt
- Input line format: [A or D] [id] [size]
 - A: Allocate, D: Deallocate
 - id: an integer identifier
 - size: bytes
- Output: size of the largest free space
 - Format: Max Free Chunk Size = \$size in bytes\$
 - Exclude the header
- We will provide you main.c and test.txt
- Your implementation must reproduce **exactly the same results** shown below

```
root@22954ec65807:/home/Lab/malloc/malloccode# LD_PRELOAD=$PWD/multilevelBF.so ./main
Max Free Chunk Size = 416
```

Grading Policy

- Produce correct answers for
 - The test.txt that TA give to you
 - Some other input files prepared by TA
- Submit your multilevelBF.c to E3
 - Filename : hw4student_ID.c
 - Example: hw4_11550999.c
- Notice It is recommended to write some testcases yourself to ensure there are no other issues. You should take care about how to free and merge the provided testcase is the simplest.

Testing OS Environment

- Ubuntu 22.04
- Install as a VM or on a physical machine

Header of your .c or .cpp

```
/*
```

```
Student No31415926
```

```
Student NameJohn Doe
```

```
Email:xxx@yyy.zzz
```

```
SE tagxnxcctxuxoxsx
```

```
Statement: I am fully aware that this program is not  
supposed to be posted to a public server, such as a  
public GitHub repository or a public web page.
```

```
*/
```

Credits

- help design this project
- Direct all questions to the current TAs